

ML ACADEMY

LLM & RAG Interview Handbook

Detailed Questions, Explanations & Worked Examples for Top Tech Interviews

Transformers • Embeddings • Retrieval • Reranking • Grounding • Evaluation RAG Architecture •
Hybrid Search • Security • Serving • Production Debugging

www.learnmlacademy.com

Free Resource - Share with your network

How to get the most from this handbook

This PDF is designed for serious interview preparation. Each topic combines an interview-ready answer with deeper explanation, examples, diagrams, code, calculations or system trade-offs wherever they improve understanding.

Recommended approach: answer the question yourself first, then compare with the explanation and work through the example.

For coding and SQL questions, state assumptions, approach, complexity and edge cases before writing code. For ML/AI system design, discuss requirements, trade-offs, failure modes, monitoring and security instead of claiming one architecture is always best.

Section	Topics
1. Research basis and further reading	1
TOTAL COVERAGE	1

1. LLM foundations you must explain clearly

Before discussing RAG, be able to explain tokenization, autoregressive generation and decoding without hand-waving.

Interview pattern Strong LLM/RAG answers localize a failure. Ask: was the problem tokenization/model behavior, retrieval, reranking, context construction, generation, authorization, evaluation or serving?

INTERVIEW QUESTION

What are tokens, embeddings and a context window?

Interview-ready answer

Tokens are the discrete units a language model processes; embeddings are learned numeric vectors representing tokens or other objects; the context window is the maximum sequence of tokens the model can condition on for a given request.

Detailed explanation

A tokenizer maps text to token IDs. The model converts those IDs to hidden vectors, adds positional information, and transforms them through layers. Retrieval embeddings are a separate use of vectors: they are designed so semantically related items are close for search. A larger context window increases how much text can be supplied, but it does not guarantee equal use of every token and increases memory/compute costs.

Example

The word 'unbelievable' may become multiple subword tokens. A 100-page policy manual can fit in some long-context models, but retrieving the 4-8 relevant sections can still be cheaper and more reliable than sending the entire document.

Likely follow-up questions

- Are token embeddings the same as sentence embeddings used for vector search?
- How does context length affect cost and KV-cache memory?
- Why do tokenizers differ across models?

Weak answer to avoid

Saying an embedding is simply a token ID converted to a float. It is a learned representation.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

Explain a decoder-only Transformer during generation.

Interview-ready answer

A decoder-only Transformer repeatedly predicts the next token using causal self-attention so each position can attend only to earlier tokens. The selected token is appended and the process repeats until a stop condition.

Detailed explanation

The input token embeddings pass through repeated blocks containing normalized self-attention, feed-forward layers and residual connections. Causal masking prevents future information leakage during training. At inference, the prompt is processed in a prefill phase, then decoding produces one or more new tokens iteratively. KV caching avoids recomputing old key/value projections.




Figure: conceptual flow for this answer.

Example

Given 'The capital of France is', the model produces a distribution over the next token. After choosing ' Paris', that token becomes part of the context for the following step.

Likely follow-up questions

- What is prefill versus decode?
- Why is decoding hard to parallelize across time?
- Where does LayerNorm appear?

Weak answer to avoid

Describing a Transformer as one attention layer; the residual, feed-forward and normalization blocks are equally important.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

Temperature and top-p: what do they do?

Interview-ready answer

Temperature rescales logits before sampling; lower values make the distribution sharper and more deterministic, while higher values flatten it. Top-p restricts sampling to the smallest set of tokens whose cumulative probability reaches p.

Detailed explanation

They are decoding controls, not measures of model confidence or factuality. Greedy or low-temperature decoding can be useful for structured tasks, but deterministic decoding does not guarantee correctness. Creative tasks may need more diversity. For tool calls and JSON, constrained or schema-guided decoding can matter more than sampling temperature.

Example

At temperature 0.2, a dominant token becomes even more dominant. With top-p=0.9, very low-probability tail tokens are excluded until the retained probability mass is about 90%.

Likely follow-up questions

- Can temperature fix hallucinations?
- What is top-k?
- What would you choose for SQL generation?

Weak answer to avoid

Saying temperature 0 means the model knows the answer with certainty.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

2. Retrieval, chunking, reranking and access control

RAG quality is mostly a systems problem. Separate recall, precision, security and context construction.

INTERVIEW QUESTION 2026 SIGNAL

Current AI system-design preparation treats RAG as a multi-stage production system rather than a vector-search demo.

Walk through an end-to-end RAG system.

Interview-ready answer

RAG retrieves external evidence at request time, places selected evidence into the model context, and asks the model to answer from that evidence. A production system separates ingestion, retrieval, reranking, authorization, generation and evaluation.

Detailed explanation

Offline, documents are parsed, chunked, enriched with metadata, embedded and indexed. Online, the query may be rewritten, filtered by user access, retrieved using lexical/vector signals, reranked, and packed into a context budget. The LLM generates a response with citations. A robust system can abstain when evidence is weak and records retrieval traces so failures can be localized.




Figure: conceptual flow for this answer.

Example

For an employee policy assistant, a vacation question retrieves only policy chunks the employee is allowed to see, reranks them, gives the top passages to the LLM, and returns an answer citing the current HR policy version.

```
def answer(question, user):
    candidates = hybrid_search(
        query=question,
        filters={"acl": user.allowed_groups},
        top_k=30,
    )
    ranked = rerank(question, candidates)[:6]
    if not enough_evidence(ranked):
        return {"answer": "Insufficient approved evidence.", "citations": []}
    return generate_grounded(question, ranked)
```

Likely follow-up questions

- Where can hallucination still happen?
- What is retrieved offline versus online?
- How do you know whether the failure was retrieval or generation?

Weak answer to avoid

Drawing only 'vector database -> LLM' and ignoring ingestion, reranking, ACLs, abstention and evaluation.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

How do you choose chunk size and overlap?

Interview-ready answer

Choose chunks large enough to preserve the evidence needed to answer a question but small enough that retrieval is precise and the context budget is not wasted. Structure-aware chunking is often better than a fixed character count.

Detailed explanation

Chunking determines what can be retrieved as one unit. Too small fragments lose definitions, tables or cross-sentence context; too large chunks dilute embedding signal and increase tokens. Overlap helps when relevant text crosses boundaries but duplicates evidence and index size. Evaluate chunking on representative queries using retrieval recall and final answer quality rather than guessing.

Example

For API documentation, chunk by heading/function with code examples kept together. For contracts, chunk by section/clauses instead of arbitrary 500-token windows.

Likely follow-up questions

- What changes for tables?
- How do parent-child retrieval or sentence windows help?
- Would you use the same chunking for keyword and vector retrieval?

Weak answer to avoid

Claiming 500 tokens with 20% overlap is universally optimal.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION 2026 SIGNAL

Modern RAG interview material increasingly expects hybrid retrieval and reranking, not vector search alone.

Vector search vs BM25 vs hybrid retrieval: when do you use each?

Interview-ready answer

BM25 is strong for exact words, identifiers and rare terms; vector search is strong for semantic similarity and paraphrases. Hybrid retrieval combines both signals and often improves recall across mixed query types.

Detailed explanation

Lexical scoring and vector similarity are on different scales, so combine them through rank fusion, calibrated score normalization or learned retrieval rather than blindly adding raw scores. Then use a reranker on the smaller candidate set. Query slices matter: error codes, names and SKUs behave differently from natural-language questions.

Example

Query 'E11000 duplicate key' benefits from lexical matching on the code. Query 'database rejects the same ID twice' may be found semantically. Hybrid retrieval can surface both before reranking.

Likely follow-up questions

- What is reciprocal rank fusion?
- Why use a cross-encoder reranker?
- How do you handle exact filters with vector search?

Weak answer to avoid

Replacing lexical retrieval just because embeddings are newer.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

What does a reranker do?

Interview-ready answer

A reranker scores a small set of retrieved candidates more precisely with the full query-document interaction, then sends only the best evidence into the context.

Detailed explanation

Initial retrieval must be cheap enough to search a large corpus, so it uses approximate lexical/vector signals. A cross-encoder or learned reranker can jointly encode the query and candidate, capturing detailed relevance but at higher cost. This two-stage architecture balances recall and precision.

Example

Retrieve 50 candidates with BM25+ANN in 30 ms, rerank the top 50 with a cross-encoder, then send the best 6 chunks to the LLM.

Likely follow-up questions

- What is the latency trade-off?
- Can an LLM be a reranker?
- How do you evaluate the reranker separately?

Weak answer to avoid

Using a very expensive reranker on the entire corpus.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION2026 SIGNAL

Enterprise RAG interviews increasingly probe identity, tenant isolation and data authorization.

How do you enforce document-level access control in RAG?

Interview-ready answer

Attach authorization metadata to indexed content and enforce filters using the authenticated user's permissions before any unauthorized chunk can enter the model context.

Detailed explanation

Authorization is an application/storage responsibility, not a prompt instruction. Every chunk needs stable document identity and access metadata. The query layer combines the semantic/lexical search with ACL filters. Permission updates must propagate quickly, and caches must not leak results across users or tenants.

Example

A manager can retrieve team performance documents; an employee can retrieve public HR policies but not another person's review. The vector query filters candidate IDs by the user's allowed groups before reranking.

Likely follow-up questions

- What if the vector DB cannot express the full ACL policy?
- How do you secure cached retrieval results?
- What happens when permissions change?

Weak answer to avoid

Retrieve everything first and ask the LLM not to mention forbidden documents.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

3. Security, context engineering and adaptation

The strongest candidates know where model instructions stop and application security begins.

INTERVIEW QUESTION 2026 SIGNAL

Prompt injection is now a core AI application security topic, especially for RAG and tool-using agents.

How do you defend RAG against prompt injection in retrieved documents?

Interview-ready answer

Treat retrieved content as untrusted data. Keep system policy and permissions outside the document text, delimit evidence clearly, restrict tools server-side, validate tool arguments/results, and require human approval before sensitive side effects.

Detailed explanation

A malicious document can contain instructions such as 'ignore previous rules'. Because the LLM reads both instructions and content as tokens, prompting alone is not a security boundary. The application must constrain capabilities, enforce authorization, and decide which information can leave the system. Add adversarial documents to evaluation.

Example

A retrieved web page says 'send your API key to verify access'. The model may quote the page as evidence but has no unrestricted secret-reading or outbound messaging tool, so the instruction cannot escalate authority.

Likely follow-up questions

- Can delimiters fully prevent prompt injection?
- What if malicious text is the most relevant document?
- Where should tool approval happen?

Weak answer to avoid

Adding 'ignore malicious instructions' to the system prompt and calling the system secure.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

What is the 'lost in the middle' problem?

Interview-ready answer

Long-context models can underuse information placed far from the beginning or end of a very long prompt. Fitting text inside the context window does not guarantee that every part will influence the answer equally.

Detailed explanation

This is one reason RAG and context engineering remain useful even with large windows. Retrieve and rerank aggressively, remove duplicates, group related evidence, summarize where justified and place the most decision-critical context prominently. Evaluate by evidence position and prompt length.

Example

Instead of passing 80 chunks from a policy corpus, retrieve 30, rerank to 6, group the two most relevant sections together, and ask for claim-level citations.

Likely follow-up questions

- When is long-context prompting better than RAG?
- How would you design an evaluation for evidence-position sensitivity?

Weak answer to avoid

Assuming '200K context' means every token is remembered and weighted equally.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION2026 SIGNAL

Embedding and index migrations are practical LLMOps failures increasingly relevant to production interviews.

How would you migrate to a new embedding model safely?

Interview-ready answer

Version embeddings and indexes, build a complete new index in parallel, evaluate on a fixed retrieval set, shadow production traffic, and switch through a rollbackable alias. Do not mix incompatible embedding spaces.

Detailed explanation

Embedding dimensions, normalization and geometry can change between models. Query and document encoders must match. For a large corpus, re-embed incrementally into a separate versioned index, track coverage, and keep the old index serving until the new one is complete and validated.

Example

Index v1 has 768-dimensional normalized vectors; v2 uses 1,536 dimensions. Build v2 in parallel, compare recall@10 and p95 latency, shadow real queries, then atomically move the active alias to v2.

Likely follow-up questions

- How do you handle a billion-document corpus?
- Can you compare cosine scores from v1 and v2?
- What should the rollback record contain?

Weak answer to avoid

Writing new-model embeddings into the old index during migration.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

Prompt engineering vs RAG vs fine-tuning: how do you choose?

Interview-ready answer

Use prompting when the model understands the task but needs clearer instructions or examples; use RAG for external, private or changing knowledge; fine-tune when you need stable behavior, style or task skill that prompting/RAG cannot reliably provide.

Detailed explanation

Fine-tuning does not automatically add fresh factual knowledge, citations or access control. RAG does not teach a new behavior if the model consistently cannot perform the task. Start with a strong baseline and diagnose the failure before adding training or infrastructure.

Example

A support assistant using outdated refund policy needs RAG. A model that must classify a specialized taxonomy consistently may benefit from fine-tuning. A model returning the wrong JSON field names may only need better schema-constrained prompting.

Likely follow-up questions

- When would you combine RAG and fine-tuning?
- What makes a fine-tuning dataset high quality?
- Can prompting fix an authorization problem?

Weak answer to avoid

Fine-tuning the model to memorize a corpus that changes every week.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

4. Evaluation, serving and production debugging

At senior levels, expect the interview to focus on evaluation boundaries, latency/memory economics and the ability to diagnose a bad answer from traces.

INTERVIEW QUESTION

How do you evaluate a RAG system end to end?

Interview-ready answer

Evaluate retrieval and generation separately, then evaluate the final product. Retrieval asks whether supporting evidence was found; generation asks whether the answer is correct and grounded in that evidence; product metrics include latency, cost, refusal behavior and user success.

Detailed explanation

Build a versioned dataset of realistic questions with expected evidence IDs or passages and, where possible, reference answers. Measure recall@k/MRR or judged retrieval relevance, then faithfulness/citation support, answer correctness and insufficiency handling. Slice by query type, corpus, user permission and freshness. Production feedback should not replace offline regression tests.

Example

If the correct chunk is absent from top-20 retrieval, generation cannot rescue the system reliably. If the chunk is present but the answer invents a number, retrieval passed and generation/grounding failed.

Likely follow-up questions

- How do you create gold evidence?
- What if multiple documents support the answer?
- How do you evaluate abstention?

Weak answer to avoid

Reporting only answer BLEU/ROUGE or only vector recall.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION2026 SIGNAL

Evaluation is now a first-class AI-engineering interview topic because production quality cannot be captured by one benchmark.

What are the strengths and weaknesses of LLM-as-a-judge?

Interview-ready answer

LLM judges scale rubric-based evaluation and pairwise comparison, but they can be biased by order, verbosity, phrasing, model family and missing domain knowledge. Calibrate them against trusted human labels and deterministic checks.

Detailed explanation

Use explicit rubrics, hide model identity, randomize answer order in pairwise tests, measure judge-human agreement, and separate factual constraints from subjective preferences. High-stakes evaluations need representative human review. Do not let the same model generate and uncritically grade its own answers.

Example

For RAG, a deterministic citation checker verifies every source ID exists, while an LLM judge scores helpfulness and whether the answer appears supported. Human reviewers audit a stratified sample of disagreements.

Likely follow-up questions

- How do you detect position bias?
- Would you use one judge or an ensemble?

- What metrics are deterministic in RAG?

Weak answer to avoid

Treating a single judge score as ground truth.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION 2026 SIGNAL

Current Hugging Face documentation exposes multiple KV cache strategies because cache memory is a real long-context serving bottleneck.

What is a KV cache, and how does continuous batching improve serving?

Interview-ready answer

KV caching reuses previous attention keys/values during autoregressive decoding; continuous batching keeps a GPU busy by dynamically admitting new requests as active sequences finish.

Detailed explanation

The cache saves repeated computation but consumes memory proportional to active context and generated tokens. Dynamic, static, offloaded or quantized cache strategies trade latency for memory. The scheduler must balance token throughput, time-to-first-token, fairness and OOM risk.

Example

If an 8-request batch has two short requests finish early, continuous batching can replace those slots with waiting requests rather than idling until all eight original sequences finish.

Likely follow-up questions

- What is prefill versus decode?
- How do long contexts affect batching?
- When would offloaded cache help?

Weak answer to avoid

Maximizing batch size without a p95 latency target or memory budget.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION

How do you make an LLM answer only from approved sources?

Interview-ready answer

Control the retrieval boundary, require citations, make insufficient-evidence behavior explicit, and verify claims against cited passages. 'Use only the context' is a helpful instruction but not sufficient enforcement.

Detailed explanation

Approved source IDs and permissions should be attached before retrieval. Context includes provenance labels. The response contract requires citations or abstention. A post-generation verifier can check that each factual claim maps to supporting evidence. High-impact domains may need human review.

Example

A compliance assistant can answer 'What is the travel reimbursement limit?' only when an approved policy passage supports the amount. If two versions conflict or no current version is found, it reports the conflict/insufficiency rather than using model memory.

Likely follow-up questions

- What if approved sources conflict?
- How do you validate a citation semantically?
- Should the model ever use pretrained knowledge?

Weak answer to avoid

Trusting the model to self-enforce the source boundary without retrieval/provenance controls.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

INTERVIEW QUESTION**How would you diagnose a RAG answer that is wrong?****Interview-ready answer**

Replay the request and inspect the trace stage by stage: query rewrite, filters, retrieved candidates, reranker ordering, context packing, model prompt, citations and final answer. Fix the earliest stage that failed.

Detailed explanation

If relevant evidence never appears, debug ingestion/chunking/index/query/retrieval. If it appears but is ranked too low, debug scoring/reranking. If good evidence is in context but the answer is unsupported, debug generation/prompt/evaluation. If the evidence is stale, fix source freshness rather than model reasoning.

Example

Question: 'What is the current parental leave duration?' Retrieval returns an old 2024 policy ranked above the 2026 policy. The correct repair is freshness/version filtering or reranking, not larger-model fine-tuning.

Likely follow-up questions

- What traces would you store?
- How do privacy constraints affect logging?
- When would you roll back the entire release?

Weak answer to avoid

Responding to every wrong RAG answer by increasing top-k or switching to a larger LLM.

Best interview habit

Answer the main question first, then volunteer one trade-off and pause for the interviewer.

Deep-dive tutorial: from first principles to production RAG

The interview questions above are useful for practice, but a strong candidate also needs a connected mental model. This section teaches the system end to end, with equations, worked examples, design choices, debugging steps and code.

How to study this section For each topic, explain the idea aloud, work the example without looking at the answer, and then connect the concept to a production failure mode. Interviews increasingly reward diagnosis and trade-off reasoning rather than memorized definitions.

1. Tokens, tokenization and why they matter in production

An LLM does not see words directly. A tokenizer converts text into a sequence of integer token IDs. Modern tokenizers typically use subword units, so a common word might be one token while a rare technical term can be split into several pieces. This matters because context limits, API billing, latency and KV-cache memory are all usually token-driven rather than character-driven.

Suppose a support request contains 4,000 English words but becomes 5,300 tokens. If the system also retrieves 8 chunks of 650 tokens each, adds a 900-token system prompt and asks for a 700-token answer, the total request can approach 12,100 tokens. A design that only counts words may unexpectedly exceed a context budget or become expensive.

Design question	Why tokenization matters
Chunk size	Chunks are usually budgeted in tokens, not characters.
Multilingual RAG	Token efficiency differs by language and tokenizer.
Latency	Longer prompts increase prefill compute and time-to-first-token.
Cost	Hosted-model pricing is commonly based on input/output tokens.
Context packing	You must reserve room for instructions, retrieved evidence and the answer.

Worked example - context budget

Assume a 16,000-token context limit. Reserve 1,000 tokens for system instructions, 700 for the user's conversation, and 1,300 for the answer. That leaves 13,000 tokens for retrieved evidence. If a chunk averages 650 tokens, the theoretical maximum is 20 chunks, but using all 20 is rarely wise because redundancy, attention dilution and latency increase. A better starting point might retrieve 30-50 candidates, rerank them, and pack only the best 5-10 chunks while measuring answer quality.

2. Self-attention: the equation behind Transformer reasoning

Inside a Transformer layer, each token representation is projected into query, key and value vectors. Scaled dot-product attention computes how strongly each query should attend to each key:

Core equation

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k}) V$$

The dot product QK^T measures compatibility. Dividing by $\sqrt{d_k}$ stabilizes the scale of logits. Softmax converts those scores into normalized weights. Multiplying by V forms a weighted combination of information from the visible tokens. In a decoder-only model, a causal mask prevents a token from attending to future positions during next-token prediction.

Tiny numerical intuition

If a query has similarity scores $[4, 1, 0]$ with three keys, softmax makes the first token dominate. If the scores are $[1.1, 1.0, 0.9]$, attention is distributed more evenly. The model learns projections so different heads specialize in different relationships; no single head should be interpreted as the whole model's reasoning.




Figure: simplified self-attention computation.

3. Embeddings for retrieval: similarity is useful, not magical

A retrieval embedding maps a query or document chunk to a dense vector. Nearby vectors are expected to represent semantically similar content. Common similarity functions include cosine similarity and dot product. Cosine similarity compares direction rather than magnitude:

Cosine similarity

$$\cos(a,b) = (a \cdot b) / (||a|| \ ||b||)$$

Embeddings are powerful for paraphrases, but they do not replace exact matching. Product IDs, error codes, legal clause numbers, model names and uncommon abbreviations often benefit from lexical retrieval. This is why production RAG frequently uses hybrid search.

Example

Query A: “database rejects the same customer ID twice” may semantically match a document about duplicate-key violations. Query B: “E11000” contains a precise error code that BM25 can match exactly. A hybrid system can retrieve both kinds of evidence, then rerank the union.

```

import numpy as np

def cosine(a, b):
    a, b = np.asarray(a), np.asarray(b)
    return float(a @ b / (np.linalg.norm(a) * np.linalg.norm(b)))

query = [0.8, 0.2, 0.1]
doc_a = [0.75, 0.25, 0.05]
doc_b = [0.1, 0.2, 0.9]
print(cosine(query, doc_a))
print(cosine(query, doc_b))
  
```

4. Approximate nearest-neighbor search and HNSW intuition

Exact nearest-neighbor search compares a query with every stored vector, which becomes expensive at large scale. Approximate nearest-neighbor indexes trade a small amount of recall for much lower latency. HNSW is a widely used graph-based approach: vectors are connected in a multi-layer proximity graph, and search navigates from coarse long-range links toward finer local neighborhoods.

Interviewers rarely expect you to derive HNSW, but they may expect you to reason about parameters such as graph construction quality, query-time search breadth, memory and recall. The practical point is that vector retrieval itself has tunable accuracy and latency; a missed document is not always an embedding-model problem.

Parameter/choice	Increasing it usually...	Trade-off
Search breadth	improves recall	higher latency
Graph connectivity	improves navigability	more memory/build cost
Candidate top-k	raises chance relevant evidence is present	more reranking/context work

5. Chunking is an information-retrieval decision

Chunking determines the unit that can be retrieved. A good chunk contains enough context to answer a question while remaining focused enough to score well for the query. Fixed-size windows are simple, but structure-aware chunking often performs better because it keeps headings, paragraphs, table rows or code blocks together.

Overlap is not free. It helps when evidence crosses a boundary but increases index size and can flood the candidate set with near-duplicates. Parent-child retrieval is another pattern: index small child chunks for precise retrieval, then return a larger parent section to the LLM for context.

Too small

High precision, but definitions and examples can be separated.

Balanced

Semantic unit stays intact with manageable token cost.

Too large

More context, but weak retrieval specificity and wasted tokens.

Experiment design

Create a labeled set of representative questions and supporting source passages. Test 300-, 600- and 1,000-token chunking, with and without overlap. Measure retrieval Recall@k first, then end-to-end answer correctness and citation quality. Choose based on the real query mix rather than a universal default.

6. BM25, dense retrieval and reciprocal rank fusion

BM25 is a lexical ranking function that rewards query-term matches while normalizing for document length and reducing the impact of very common terms. Dense retrieval uses embeddings. Hybrid retrieval combines candidate lists so the system can handle both exact-term and semantic queries.

A simple robust fusion method is Reciprocal Rank Fusion (RRF). Instead of combining incomparable raw BM25 and vector scores, RRF gives each document a score based on rank:

RRF

$$\text{score}(d) = \sum 1 / (k + \text{rank}_i(d))$$

where k is a smoothing constant and $\text{rank}_i(d)$ is the document's rank in retrieval list i . The exact constant should be validated for the system; the key idea is that high ranks across multiple retrievers accumulate evidence.

```
def rrf(rankings, k=60):
    scores = {}
    for ranking in rankings:
        for rank, doc_id in enumerate(ranking, start=1):
            scores[doc_id] = scores.get(doc_id, 0.0) + 1.0 / (k + rank)
    return sorted(scores, key=scores.get, reverse=True)

lexical = ["d7", "d2", "d9"]
vector = ["d2", "d4", "d7"]
print(rrf([lexical, vector]))
```

7. Reranking: spend expensive computation only where it helps

The first retrieval stage must search a large corpus quickly, so it uses cheaper signals. A reranker evaluates a much smaller candidate set more precisely. Cross-encoder rerankers jointly process the query and each candidate and can capture token-level interactions that independent embeddings miss.

A common architecture is retrieve 30-100 candidates, rerank them, then pass perhaps 4-10 diverse evidence chunks to the LLM. Larger candidate sets can improve recall but increase reranking latency. Measure the quality/latency curve rather than assuming more is always better.

Latency budget

Suppose retrieval takes 35 ms, reranking 50 candidates takes 90 ms, context assembly 10 ms, and model generation dominates at 1.4 s. If the product target is p95 under 2 s, doubling reranking work may still fit. If the product target is 500 ms, you may need a smaller candidate set, a faster reranker or asynchronous architecture.

8. Context construction: relevance, diversity and position

Even with good retrieval, the system can fail during context assembly. Duplicated chunks waste tokens, contradictory versions confuse the model, and important evidence can be buried. Context construction should consider relevance, diversity, freshness, source authority, permissions and token budget.

Maximal Marginal Relevance (MMR) is one way to trade relevance against redundancy. It favors candidates that are relevant to the query but different from already selected evidence. Another approach is rule-based deduplication using document IDs, section IDs or embedding similarity.

Production habit Preserve provenance. Every context block should carry a stable source/document/chunk identifier, version or timestamp when relevant, and permission metadata. This makes citations, debugging and freshness checks possible.

9. Grounding, citations and abstention

“Answer only from the context” is an instruction, not a security or correctness guarantee. A grounded RAG system defines an explicit response contract: cite supporting evidence, abstain when evidence is insufficient, and surface conflicts when approved sources disagree.

Groundedness can be evaluated at claim level. Split the answer into factual claims, retrieve the cited passages, and check whether each claim is entailed by the evidence. Some checks can be deterministic (citation ID exists, quoted number is present); semantic support often needs a classifier, LLM judge or human review.

```
def validate_response(answer, citations, allowed_ids):
    if not set(citations).issubset(allowed_ids):
        return False, "unknown citation"
    if not citations and contains_factual_claim(answer):
        return False, "factual answer without evidence"
    return True, "basic provenance checks passed"
```

10. Evaluate retrieval separately from generation

An end-to-end “answer accuracy” number cannot tell you what broke. Maintain stage-specific metrics. Retrieval Recall@k asks whether at least one relevant item appeared in the top k. Precision@k measures how much of the retrieved set is relevant. MRR rewards systems that place the first relevant item high. NDCG is useful when there are graded relevance labels.

Stage	Useful metrics	Typical question
Retrieval	Recall@k, Precision@k, MRR, NDCG	Did we find the right evidence?
Reranking	NDCG, MRR, pairwise win rate	Did we move the best evidence upward?
Generation	correctness, groundedness, citation accuracy, refusal quality	Did the answer use evidence correctly?
Operations	p50/p95 latency, token cost, cache hit rate, error rate	Is the system reliable and affordable?

Recall@k example
For 100 evaluation questions, suppose the gold supporting document is present in the top-5 retrieval results for 87 questions. Recall@5 = 87/100 = 0.87. If final answer accuracy is only 0.61, retrieval is not the only problem; investigate reranking, context construction and generation.

11. LLM-as-a-judge: useful, but calibrate it

LLM judges can score helpfulness, groundedness, style or pairwise preference at scale. They are especially useful during rapid iteration, but they can exhibit position bias, verbosity bias, sensitivity to rubric wording and model-family effects. High-quality evaluation uses explicit rubrics, blind comparisons, randomized order and a human-calibrated sample.

Do not ask one model a vague question like “Is this answer good?” Instead define dimensions and anchors. For groundedness, the judge should see the question, answer and evidence and decide whether each claim is supported. For safety, use adversarial test suites and deterministic policy checks where possible.

12. Fine-tuning vs RAG vs long context


Need	RAG	Fine-tuning	Long-context prompting
Fresh private knowledge	Strong fit	Weak fit alone	Possible for small corpora
Change style/format/behavior	Limited	Strong fit	Prompt-dependent
Source citations	Natural fit	Difficult	Possible if sources supplied
Frequent knowledge updates	Easy to re-index	Requires new training	Supply latest context
Very large corpus	Efficient selection	Not a database	May be too expensive

These are not mutually exclusive. A production system can fine-tune a model for response style or task behavior while using RAG for fresh domain knowledge. Long context can reduce retrieval complexity for small bounded corpora, but it does not automatically solve permissions, freshness, provenance or equal attention to all supplied information.

13. Security: prompt injection is a capability problem

Retrieved documents, web pages and tool outputs are untrusted data. A malicious document can contain instructions that try to override the application's intent. Because the model consumes instructions and data in the same token stream, delimiters and prompt wording reduce risk but do not create a hard security boundary.

Defenses belong at multiple layers: least-privilege tools, server-side authorization, input/source trust policies, output validation, sandboxing, human approval for high-impact actions, secret isolation, audit logs and adversarial evaluation. For RAG specifically, permissions should filter documents before their content reaches the model.



14. Multi-tenant RAG and access-control design

For enterprise systems, tenant isolation and document-level access control are core architecture concerns. Every indexed chunk should preserve stable ownership and ACL metadata. Query execution must combine relevance retrieval with the authenticated user's allowed scope. Caches require the same care: a cached response or candidate set for one tenant must never be reused for another tenant unless the cache key and data are provably safe.

Failure scenario

A vector index stores HR policies and employee reviews. If the application retrieves top-k globally and then asks the LLM to ignore unauthorized chunks, the sensitive text has already crossed the security boundary. Correct design filters the searchable candidate space before reranking/context assembly.

15. Serving LLMs: prefill, decode and KV-cache memory

Generation has two distinct phases. **Prefill** processes the input prompt and can use parallel matrix operations across many input tokens. **Decode** generates new tokens one at a time, so each sequence has a serial dependency. KV caching stores past key/value states so old projections do not need to be recomputed every step.

KV cache reduces compute but consumes memory that grows with active sequence length, layer count, number of KV heads and precision. Long-context requests can therefore reduce how many concurrent users fit on a GPU. Modern serving systems use strategies such as paged/block KV memory, prefix caching, cache quantization/offloading, chunked prefill and continuous batching.

Operational intuition

A 12k-token prompt may have a slow time-to-first-token because prefill is large. After prefill, token generation can still be fast per token. A short prompt that asks for a very long answer has the opposite shape: small prefill, long decode. Capacity planning must account for both.

16. Continuous batching and scheduling

Static batching waits for a fixed set of requests to finish together. This wastes GPU capacity when sequences have different lengths. Continuous batching admits new requests as old sequences finish, improving utilization and throughput. The scheduler must balance throughput with fairness and latency objectives.

Metric	Meaning
TTFT	Time to first token; sensitive to queueing and prefill.
TPOT	Time per output token during decode.
Throughput	Tokens or requests completed per unit time.
p95 latency	Tail experience for slower requests.

17. Production debugging playbook

When a RAG answer is wrong, inspect the trace in causal order. The earliest incorrect stage is usually where the fix belongs.

- 1. **Ingestion:** Was the correct source parsed, versioned and indexed?
- 2. **Chunking:** Is the needed evidence preserved in a retrievable unit?
- 3. **Query processing:** Did rewriting change the intent?
- 4. **Authorization:** Were valid sources filtered out or forbidden sources allowed?
- 5. **Retrieval:** Was relevant evidence in the candidate set?
- 6. **Reranking:** Was it promoted high enough?
- 7. **Context packing:** Was it included without being truncated?
- 8. **Generation:** Did the model use the evidence correctly?
- 9. **Validation:** Did citations/guardrails catch unsupported claims?


```
trace = {
  "query": question,
  "rewrite": rewritten_query,
  "filters": acl_filters,
  "retrieved": [(d.id, d.score) for d in candidates],
  "reranked": [(d.id, d.score) for d in ranked],
  "context_ids": [d.id for d in packed],
  "model": model_name,
  "latency_ms": latency,
}
# Store only what policy/privacy rules permit.
```

Two complete system-design walkthroughs

Case study A - enterprise policy assistant

Goal: employees ask questions about HR, travel and finance policies and receive a current answer with citations. The system must prevent cross-tenant or role-based information leakage.

1. **Requirements:** p95 response under a few seconds, citations required, current version preferred, abstain on missing/conflicting evidence.
2. **Ingestion:** parse approved document stores, retain document ID, section hierarchy, effective date, owner, ACL and checksum.
3. **Chunking:** split by heading/clauses; keep tables and their headers together; use parent-child retrieval for long sections.
4. **Indexing:** lexical + dense indexes; update changed documents incrementally and remove superseded versions.
5. **Online query:** authenticate user, derive ACL filters, optionally rewrite query, execute hybrid retrieval, fuse ranks, rerank top candidates.
6. **Context:** remove duplicates, prefer current authoritative sources, fit evidence into a token budget with stable citation IDs.
7. **Generation:** answer from supplied evidence, cite every important factual claim, abstain when evidence is insufficient.
8. **Validation:** verify citation IDs, detect unsupported numeric claims, optionally run a groundedness judge.
9. **Observability:** store safe traces, retrieval metrics, latency breakdowns and user feedback.



What can go wrong?

- Old policy outranks new policy → add version/freshness logic.
- User cannot find authorized document → inspect ACL propagation and filters.
- Correct passage retrieved but not used → inspect context position, prompt and generation.
- Citations look plausible but do not support claim → add semantic citation validation.
- Latency spikes → break down retrieval, reranking, prefill and decode time.

Case study B - customer-support troubleshooting assistant

Goal: answer product questions from manuals, runbooks and resolved tickets while avoiding confidential customer-data leakage.

The query mix contains exact error codes and natural-language descriptions, so hybrid retrieval is especially valuable. Tickets need aggressive PII redaction and access filtering. Freshness matters because product versions change. The system should retrieve version-compatible documentation, cite the relevant procedure, and escalate when confidence is low.

Interview trade-off Resolved tickets can contain highly useful long-tail fixes but are noisy and may include bad workarounds. Treat them as lower-authority evidence than official runbooks, or use source-quality features during reranking.

Additional interview drill: 20 questions to answer aloud

#	Question	What a strong answer must include
1	Why can a larger context window still benefit from RAG?	Cost, latency, provenance, permissions, freshness, attention dilution.
2	When would BM25 outperform embeddings?	Exact identifiers, rare terms, codes, names and lexical constraints.
3	What is Recall@k?	Fraction of queries where relevant evidence appears within top k.
4	Why rerank after retrieval?	Cheap high-recall stage followed by precise expensive ranking.
5	How do you handle tables in RAG?	Structure-aware parsing, preserve headers, row/group retrieval or table-specific representation.
6	How do you handle stale documents?	Version metadata, effective dates, source-of-truth policy, incremental re-indexing.
7	What is query rewriting?	Transform query for retrieval while preserving intent; evaluate because rewrite can harm.
8	What is multi-query retrieval?	Generate complementary query variants to improve recall, then fuse/dedupe.
9	How would you evaluate chunk size?	Labeled queries, retrieval metrics and end-to-end answer quality.
10	How do you prevent cross-user cache leaks?	Permission-aware cache keys, isolation, avoid caching private context unsafely.
11	What is prompt injection?	Untrusted input tries to redirect model behavior; control capabilities outside prompts.
12	Fine-tuning or RAG for changing company policies?	RAG for freshness/provenance; fine-tuning may complement behavior/style.
13	What is TTFT?	Time to first token; queueing + prefill are major contributors.
14	Why does KV cache use memory?	Stores per-layer past key/value tensors growing with sequence length.
15	Why continuous batching?	Replace completed sequences dynamically to improve GPU utilization.
16	What does abstention mean?	System declines to answer when approved evidence is insufficient or conflicting.
17	How do you test prompt-injection defenses?	Adversarial documents/queries, capability checks, tool/ACL enforcement and red-team evals.
18	How do you debug low Recall@5?	Inspect parsing, chunks, embedding model, lexical channel, filters and ANN settings.
19	What is citation accuracy?	Whether cited evidence actually supports the associated claims.
20	What makes a production RAG answer strong in an interview?	Architecture + failure modes + metrics + security + latency/cost trade-offs.

Hands-on mini lab: build and evaluate a tiny hybrid RAG pipeline

The following code is intentionally small enough to discuss in an interview. A production system would use robust tokenization, persistent indexes, proper authentication, observability and tested model clients.

```
from dataclasses import dataclass
from collections import Counter
import math

@dataclass
class Doc:
    id: str
    text: str
    groups: set

DOCS = [
    Doc("travel-2026", "Travel reimbursement limit is 150 dollars per day.", {"all"}),
    Doc("travel-2024", "Travel reimbursement limit is 120 dollars per day.", {"all"}),
    Doc("mgr-review", "Manager-only performance review guidance.", {"managers"}),
]

def lexical_score(query, text):
    q = Counter(query.lower().split())
    d = Counter(text.lower().split())
    return sum(min(q[t], d[t]) for t in q)

def retrieve(query, user_groups, k=5):
    allowed = [d for d in DOCS if d.groups & user_groups]
    ranked = sorted(allowed, key=lambda d: lexical_score(query, d.text), reverse=True)
    return ranked[:k]

print([d.id for d in retrieve("travel reimbursement limit", {"all"})])
```

Exercises:

1. Add a freshness boost so the 2026 policy outranks the 2024 policy.
2. Add a toy dense-similarity score and combine the two ranked lists with RRF.
3. Create five labeled queries and compute Recall@2.
4. Add an authorization test proving a non-manager never receives mgr-review.
5. Add an abstention rule when top evidence is weak or conflicting.

Interview lesson

A small system you can explain, test and debug is more convincing than listing five vector databases and three frameworks without understanding the retrieval/evaluation loop.

Hands-on labs and worked exercises

These exercises turn the concepts into things you can calculate, code, debug or design. Work them before reading the expected reasoning.

Lab 1 - Calculate Recall@k

Goal Measure retrieval before blaming the LLM.

Steps

1. Create 10 queries with one known supporting document each.
2. Retrieve top-3 document IDs.
3. Count queries where the gold document appears.
4. Compute Recall@3.

Expected reasoning / result

If 8/10 queries contain the gold document, Recall@3=0.8. Generation quality cannot recover evidence that was never retrieved.

Interview connection

Separate retrieval evaluation from answer evaluation.

Lab 2 - Hybrid retrieval with RRF

Goal Fuse lexical and dense rankings without adding incomparable raw scores.

Steps

1. Take lexical ranking [A,B,C,D] and dense ranking [C,A,E,B].
2. Use RRF with k=60.
3. Compute scores and sort.
4. Explain why documents high in both lists rise.

```
def rrf(lists,k=60):
    s={}
    for ranking in lists:
        for r,d in enumerate(ranking,1):
            s[d]=s.get(d,0)+1/(k+r)
    return sorted(s.items(), key=lambda x:x[1], reverse=True)

print(rrf([["A","B","C","D"],["C","A","E","B"]]))
```

Expected reasoning / result

RRF uses rank positions rather than raw BM25/vector scores, making fusion simple and robust.

Interview connection

Then discuss reranking the fused candidate set.

Lab 3 - Permission test

Goal Prove unauthorized chunks never reach context.

Steps

1. Create public, manager-only and finance-only chunks.
2. Run retrieval as three user roles.
3. Assert forbidden document IDs are absent before reranking and generation.
4. Repeat after permission changes and cache hits.

Expected reasoning / result

Authorization must constrain retrieval/candidate access. A prompt telling the model “do not reveal private text” is not sufficient.

Interview connection

Security tests should be deterministic wherever possible.

Lab 4 - Diagnose stale-answer failure

Goal Trace a wrong answer to the earliest failing stage.

Steps

1. Question asks current travel limit.
2. Retrieved top results contain 2024 and 2026 policy, but 2024 ranks first.
3. Check version metadata, freshness filters and reranker features.
4. Fix retrieval/ranking rather than fine-tuning the LLM.

Expected reasoning / result

The model received misleading evidence; the root cause is retrieval freshness/ranking.

Interview connection

Strong answers repair the earliest broken stage.

Lab 5 - Token and latency budget

Goal Design context packing under constraints.

Steps

1. Context limit=32k. Reserve 2k instructions/chat and 2k answer.
2. Candidate chunks average 900 tokens.
3. Choose a maximum packed count and explain why you leave margin.
4. Set TTFT and total latency metrics.

Expected reasoning / result

The theoretical maximum is ~31 chunks, but far fewer may be optimal because redundancy, prefill cost and evidence dilution grow. Measure quality vs tokens.

Interview connection

Context-window size is not an instruction to fill it.

Last-minute revision sheet

Use this page in the final 30-60 minutes before an interview. Say each answer aloud in one or two sentences; if you cannot, revisit the relevant chapter.

Topic	One-line recall
RAG flow	Ingest/index -> authorized hybrid retrieval -> rerank -> context -> generate -> validate/cite -> answer or abstain.
Chunking	Preserve semantic units; evaluate chunk size/overlap on real queries instead of copying defaults.
Hybrid retrieval	Lexical exact-match strength + vector semantic strength, then rerank.
ACLs	Filter before unauthorized text reaches the LLM; never rely on prompt-only access control.
Prompt injection	Retrieved content is untrusted data; application code enforces capabilities and approvals.
Evaluation	Separate retrieval recall, answer correctness, groundedness/citations, refusal, latency and cost.
LLM judge	Useful scalable rubric, but calibrate against humans and deterministic checks.
Serving	KV cache saves compute; continuous batching improves GPU utilization but consumes memory and affects latency.

Research basis and further reading

This handbook is original LearnMLAcademy material. Public interview reports and official technical documentation were used to identify current topic emphasis and production practices. The explanations, examples, diagrams and code in this handbook are independently written.

Hugging Face Transformers - KV Cache Strategies

https://huggingface.co/docs/transformers/kv_cache

Current cache types, trade-offs, offloading, quantized cache and iterative generation.

Aced / Exponent - System Design Interview Prep (2026)

<https://www.tryexponent.com/blog/system-design-interview-guide>

Current system-design prep highlights cost, failure modes, operational judgment, AI/LLM infrastructure and agentic systems.

InterviewQuery - Google Data Scientist Interview Guide (2026)

<https://www.interviewquery.com/guides/google-data-scientist>

Recent candidate reports emphasize SQL window functions, A/B testing, product metrics, ML fundamentals and behavioral rounds.

Aced / Exponent - Meta Senior MLE Interview Experience (2026)

<https://www.tryexponent.com/experiences/meta-facebook-senior-machine-learning-engineer-interview-8719a7>

Candidate-reported Meta loop covering classic coding, AI-assisted debugging, behavioral questions and recommendation-system design with online/offline evaluation.

Google Machine Learning Crash Course - Production ML Systems

<https://developers.google.com/machine-learning/crash-course/production-ml-systems>

Useful production-systems framing: the model is only one part of a larger data, serving, testing and monitoring ecosystem.

OpenAI Cookbook - Evals

<https://developers.openai.com/cookbook/topic/evals>

Current evaluation recipes covering model and agent quality, traces, security and iterative improvement.

Anthropic - Trustworthy Agents in Practice (2026)

<https://www.anthropic.com/research/trustworthy-agents>

Current discussion of agent autonomy, unintended actions and prompt-injection risks.

Keep learning at learnmlacademy.com

167 in-depth lessons • Python code examples • Visual diagrams • Projects • Interview preparation • Generative AI • LLM & RAG • Agentic AI • Always free